

COS470 Final Study Guide

Introduction & Agents

- What is AI?
 - o Getting computers to do intelligent things
 - o Medical diagnostic reasoning, planning, designing... doing the right thing and being rational
- What is AI technique?
 - o The method of exploiting knowledge through search, structure of knowledge and abstraction
- **Tractability** – a problem is called intractable if the time required to solve instances of the problem grows exponentially with the size of the instances.
- Ideal agent: responsive, proactive, rational, maximizes utility, autonomous, social
- Properties of Agents: sensors, actuators, knowledge of goal, knowledge of utility, can it change via learning
- Good Agent: rational, does the right thing with the information it has
 - o Should be in relation to human goals
 - o Some agents know their performance measure
 - o Best performance measure: relate to outcome wanted, not how agent behaves
- Rational Agent: For each possible sequence, a rational agent should select an action that is expected to maximize its performance measure, given the evidence provided by the percept sequence and whatever built in knowledge the agent has
- Task Environments
 - o Properties of world
 - o Circumscribed to problem and things affecting solution
 - o Domain impacts task environment properties
 - o Complexity of environment -> effectively nondeterministic
 - o Sensor/actuator noise -> effectively nondeterministic
 - Deterministic
 - Static
 - Episodic
 - Sequential
 - Discreet
 - Continuous
- Types of problems:
 - o Classification/Analysis
 - Image recognition
 - Diagnosis
 - Data mining
 - NLP
 - Sentiment analysis
 - o Synthesis Construction
 - Automated planning
 - Scheduling
 - Designing buildings, car parts, new drugs

- Divide and conquer method for decomposable problems into subproblems”
 - P is $O(n^2) \rightarrow O(m \cdot 2^{(n/m)})$
- Types of agents
 - Reflex Agent – simplest has condition action rules that matches what the world is like now, atemporal, no history (condition action pairs)
 - Model Based Agent – Consider what we’ve done in the past, How the world evolves over time. How actions affect the world
 - Goal based agent – what would happen if I did this, compare alternatives, plan out ahead, search agent
 - Utility based agent – what do I want to happen
 - Learning Agent – learning, how to do things better by how well its doing something.

Search Space

- State-space search
 - Modeled with graph, usually digraph
 - Problem = state space + initial state + description of final state
 - State: the configuration of the world
 - State space: collection of all possible states
 - Transitions or links: events or processes in the world, Agent’s, or other Agent’s actions
 - Solution: path from start to goal states.
- Uninformed search
 - Complete? optimal?
 - BFS – complete, optimal as long as cost of edges are equal
 - DFS – complete, not optimal
 - IDDFS – complete, optimal
- Heuristic search & heuristics
 - Use knowledge to prioritize nodes
 - Heuristic rule: search space topology, problem domain property
 - Heuristic function maps state (local information how good, how good are the next states) (Global information, how close is this state/next state compared to goal)
 - A* - complete and optimal
- Local & online search
 - Hillclimb – local best first search
 - Not optimal, not complete
 - Backtracking HC – continuing if there exists child s' and $h(s') > h(s)$
 - Otherwise backtrack to previous choice point
 - Can check for repeated states
 - If can recognize goal and it’s a local minimum, backtrack.
 - Simulated annealing
 - **Not optimal, complete**
 - Instead of random jumps, try suboptimal states
 - Start – probability of random moves is high (decreases over time/jumps)
 - At node take a random move
 - if better take it
 - if worse – take it some percentage of the time
 - Beam search
 - BFS branching factor is high

- If we can reduce b that's great!
 - Search only a certain beam width (maybe **not optimal, might not be complete**)
 - Stochastic Beam Search
 - Like beam search with random elements
 - Choose l nodes at random
- Constraint-satisfaction
 - Constraint Graph
 - Nodes = variables
 - Arcs = constraints
 - Can treat CSP as search problem
 - Which variable to set
 - Pick variable with smallest remaining domain
 - Reduces branching factor = fewest alternatives to backtrack to
- Adversarial search
 - Minimax and alpha beta pruning
 - Making sure your opponent doesn't get you to make a bad move
 - Alpha(max)-beta(beta) cutoffs $O(b^d)$ worst case
 - Best case $O(b^{d/2})$ if nodes in worst case order

Neural networks: Searching to create agents

- Search and ANNs
 - Search comes in during training – ANNs learn by adjusting weights or parameters
 - Goal: find the best value for each weight
- Neurons – a thing that holds data (activation)
 - Output number = confidence
 - Activations from one layer determine activations passed to the next layer (hidden/output)
- Training
 - Activation sequence is known through experimentation
 - Weights are set and modified with training
 - Why layers?
 - Components are processed individually
 - Activation weight = weighted sum
 - We want $0 \rightarrow 1$
 - Use activation function to force weight to $0 \rightarrow 1$
 - sigmoid function
 - rectified linear unit (easier to compute no exponentiation)
 - bias
 - force neuron to activate only if the weighted sum $>$ the bias
 - so $w' = w - \text{bias}(\alpha)$
- Learning
 - Just find the right weights and biases
 - Gradient decent (ball rolling down a hill) gradient of the loss function (how far off we are) update weights to minimize loss.
 - Backpropagation is the algorithm that computes the gradient.

- Weights are modified along the steepest descent of the gradient to minimize error efficiently

Types of Neural Networks

- Types of neural nets
 - Convolutional
 - Great for image recognition
 - LSTM
 - Natural language processing, time series prediction
 - Multilayer perceptron
 - Basic NN

Knowledge Representation

- Five types of knowledge
 - Declarative knowledge – concepts, facts objects
 - Structural knowledge – problem solving knowledge, relationship between concepts and objects
 - Procedural knowledge – know how to do something, rules, strategies, procedures
 - Meta-knowledge – knowledge about knowledge
 - Heuristic Knowledge – expert knowledge in field
- Cycle of knowledge representation
- Knowledge versus intelligence
 - Knowledge fuels intelligence
 - Knowledge is the base of intelligent behavior
- Techniques
 - Logical Representation
 - Language with definite rules dealing with propositions
 - No ambiguity
 - Concludes on conditions and cements communication rules
 - Syntax (how we construct in logic define symbols) and semantics (rules of governing interpretation of sentences, prescribes meaning)
 - Logical reasoning, basis of programming languages
 - Semantic Network Representation
 - Alternative of predicate logic
 - Graphical networks
 - Nodes are objects, edges are relationships
 - Isa, kindof, etc
 - Natural representation... simple, more compute time, not intelligent
 - Frame Representation
 - Collection of attributes describing an entity in the world
 - Slots, facets, groups related data, easily understandable, generalization
 - Production Rules
 - Check if condition exist
 - If yes, production rule fires and action is executed
 - 3 parts:

- Set of rules
- Working memory
- Recognize act cycle
- Natural language with no learning capabilities & inefficient
-
-

Planning and acting

Backward Chaining

Work backward from goal, what rules can conclude something about hypothesis

Needs to account for uncertainty (probability), incomplete (Bayesian Nets)

Forward Chaining

- Automated planning
- Handling uncertainty
- Decision making with utilities

Machine learning

- Symbolic ML
- Going beyond MLPs (Assignment 6: Convolutional Neural Networks)
- Reinforcement learning


Interacting with others

- Natural language processing
- Multiagent systems



A comparison table between DFS, BFS and IDDFS

	Time Complexity	Space Complexity	When to Use ?
DFS	$O(b^d)$	$O(d)$	=> Don't care if the answer is closest to the starting vertex/root. => When graph/tree is not very big/infinite.
BFS	$O(b^d)$	$O(b^d)$	=> When space is not an issue => When we do care/want the closest answer to the root.
IDDFS	$O(b^d)$	$O(bd)$	=> You want a BFS, you don't have enough memory, and somewhat slower performance is accepted. In short, you want a BFS + DFS.


Minimax Alpha beta pruning



Simple Neural network



A sigmoid unit in a neural network



Domain: $(-\infty, +\infty)$

Range: $(0, +1)$

$\sigma(0) = 0.5$

Other properties

$$\sigma(x) = 1 - \sigma(-x)$$


$$\sigma(x) = \frac{1}{1 + e^{-x}} = \frac{e^x}{e^x + 1}$$

$$\sigma'(x) = \sigma(x)(1 - \sigma(x))$$

Inference by enumeration

$$P(b, j, m) = \sum_e \sum_a P(b)P(e)P(a|b, e)P(j|a)P(m|a)$$

$$= P(b) \sum_e \sum_a P(e)P(a|b, e)P(j|a)P(m|a)$$



$$\begin{aligned}
 &= 0.001 \times (0.001 \times 0.95 \times 0.9 \times 0.7 + \quad e, a \\
 &\quad 0.001 \times 0.05 \times 0.05 \times 0.01 + \quad e, \neg a \\
 &\quad 0.999 \times 0.94 \times 0.9 \times 0.7 + \quad \neg e, a \\
 &\quad 0.999 \times 0.06 \times 0.05 \times 0.01) \quad \neg e, \neg a
 \end{aligned}$$